

Retail Media Monetization Benchmark, India - Recommendation Memo

Why this project

Built for the Tata CLiQ Monetization Intern application (retail-media reporting and brand-pitch market research) and directly relevant to equity research roles, where ad monetization is a live theme for Indian internet-platform coverage. The question: how much ad revenue does an Indian marketplace extract per rupee of GOV, benchmarked against comparable listed companies.

Methodology

1. **Source three listed marketplaces' own disclosures** - annual reports, investor presentations, shareholder letters, and (where primary PDFs didn't parse) reputable financial-press summaries of the same filings.
2. **Tag every disclosed figure by** `metric_type` in `data/raw/marketplace_disclosures.csv` - `gov_gmv`, `ad_revenue_pct_of_gov`, `ad_revenue_pct_of_gov_TARGET` (forward-looking, not current), `ad_SPEND_not_ad_revenue_EXCLUDED`, `bundled_revenue_not_ad_only`, `blended_take_rate_not_ad_only`. This forces the ad-spend/ad-revenue distinction into the data itself, not just into prose that's easy to skim past.
3. **Compute an implied ad-revenue figure only where both a GOV figure and a standalone ad-revenue-% figure exist for the same company/segment** (`src/benchmark.py`). Every other case is written as `INSUFFICIENT_DISCLOSURE`, never estimated.

Results

Swiggy is the only company with a standalone ad-revenue-as-%-of-GOV disclosure: $\geq 4\%$ in food delivery (a disclosed floor - the real figure could be higher), and a forward target of 6-7% in quick commerce (explicitly described as an expectation, not a current achieved figure - used at the 6.5% midpoint and labelled as a target throughout). Applied to real Q4 FY26 GOV (Food Delivery INR 9,005cr, Instamart INR 7,881cr), this implies $\geq$ INR 360.2cr and INR 512.3cr respectively.

Nykaa's GMV is well disclosed (FY25 INR 15,604cr, FY26 INR 19,963cr) but no standalone ad-revenue line was found - only its own INR 995cr FY25 marketing/ad *spend* figure, which is the wrong side of the transaction and was excluded. Eternal's revenue is disclosed (Q4 FY26 INR 17,292cr, Blinkit INR 13,232cr) but not a GOV/GMV figure comparable to the other two - and Blinkit's move to an inventory-led accounting model from Q1 FY26 makes its revenue figure not directly comparable to a GOV-based metric without further adjustment not attempted here.

What broke

Two primary-source PDFs (Nykaa's Q4 FY26 investor presentation, Eternal's Q4 FY26 shareholder letter) were fetched but didn't parse: one returned only its digital-signature certificate metadata, the other exceeded the fetch tool's size limit. Rather than guess at their contents or fall back to unreliable numbers, this project relies on financial-press summaries of the same underlying filings and documents the limitation plainly - see `LIMITATIONS.md`.

The more consequential "what broke": an early read of the search results nearly used Nykaa's and Eternal's disclosed "AdEx"/ad-spend figures as if they were ad-revenue figures, which would have produced a clean-looking but wrong three-company table. Re-reading the source context (these are figures about money the companies spend advertising themselves, not revenue they earn from brands) caught the error before it reached the data file.

What this demonstrates for the role

The habit the JD actually tests for - read a company's own disclosure carefully enough to know what it does and doesn't say, rather than pattern-matching a plausible-looking number into a benchmark table. This is the same discipline as the RTO revenue-basis catch and the RBI ECL reconciliation elsewhere in the portfolio.